

Car price prediction

In this notebook we are going to do a regression project to predict selling price of a car based on some features that are provided in dataset.

Data Dictionary

- **Car_name** -----> The name of the car.
- **Year** -----> Year of car production.
- **Selling_price** -----> The price that car was dealt with: target value
- **Present_price** -----> Current car's price .
- **Kms_Driven** -----> The distance traveled by the car in kilometer.
- **Fuel_Type** -----> Fuel type that car uses.
- **Seller_type** -----> Car has sold be a dealer or a individual.
- **Transmission** -----> Type of gearbox that car has.
- **Owner** -----> The number of previous owner.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 # Defining the path
        2 path = 'E:/EDU/Programming/Python/AI/car_prediction_data.csv'
```

```
In [3]: 1 dataset = pd.read_csv(path)
        2 dataset.head()
```

```
Out[3]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manua
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manua
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manua
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manua
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manua

Data preprocessing and Exploratory data analysis.

```
In [4]: 1 # see the unique values
        2 print(dataset['Fuel_Type'].unique())
        3 print(dataset['Seller_Type'].unique())
        4 print(dataset['Transmission'].unique())
        5 print(dataset['Owner'].unique())
```

```
['Petrol' 'Diesel' 'CNG']
['Dealer' 'Individual']
['Manual' 'Automatic']
[0 1 3]
```

```
In [5]: 1 # finding null values
        2 dataset.isnull().sum()
```

```
Out[5]: Car_Name      0
        Year          0
        Selling_Price  0
        Present_Price  0
        Kms_Driven     0
        Fuel_Type      0
        Seller_Type     0
        Transmission   0
        Owner          0
        dtype: int64
```

```
In [6]: 1 # Car name is not necessary as it has no effect on price of the car
        2 dataset = dataset.drop('Car_Name', axis = 1)
```

```
In [7]: 1 dataset.head()
```

```
Out[7]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

```
In [11]: 1 # insted of year of the car production it is better to have the age of it. us
        2 import datetime
        3 current_time = datetime.datetime.now()
        4 current_year = current_time.year
        5 dataset['Car_age'] = current_year - dataset['Year']
```

In [12]:

```
1 dataset.head()
```

Out[12]:

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	C
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0	
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0	
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0	
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0	

In [13]:

```
1 # know we do not need Year column
2 dataset = dataset.drop('Year', axis = 1)
3 dataset.head()
```


Out[13]:

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Car_age
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	9
1	4.75	9.54	43000	Diesel	Dealer	Manual	0	10
2	7.25	9.85	6900	Petrol	Dealer	Manual	0	6
3	2.85	4.15	5200	Petrol	Dealer	Manual	0	12
4	4.60	6.87	42450	Diesel	Dealer	Manual	0	9

Let's see the relation between independent variables and dependent one which is 'Selling_Price'

Selling_Price VS Present_Price


```
In [14]: 1 sns.scatterplot(x = 'Present_Price', y = 'Selling_Price', data = dataset).se
```



The scatter plot displays the rise in selling price in relation to the current price. Nonetheless, there are certain cars whose price is comparatively lower than others with a similar current price. Hence, there must be additional factors influencing the selling price.

Selling_Price VS Kms_Driven

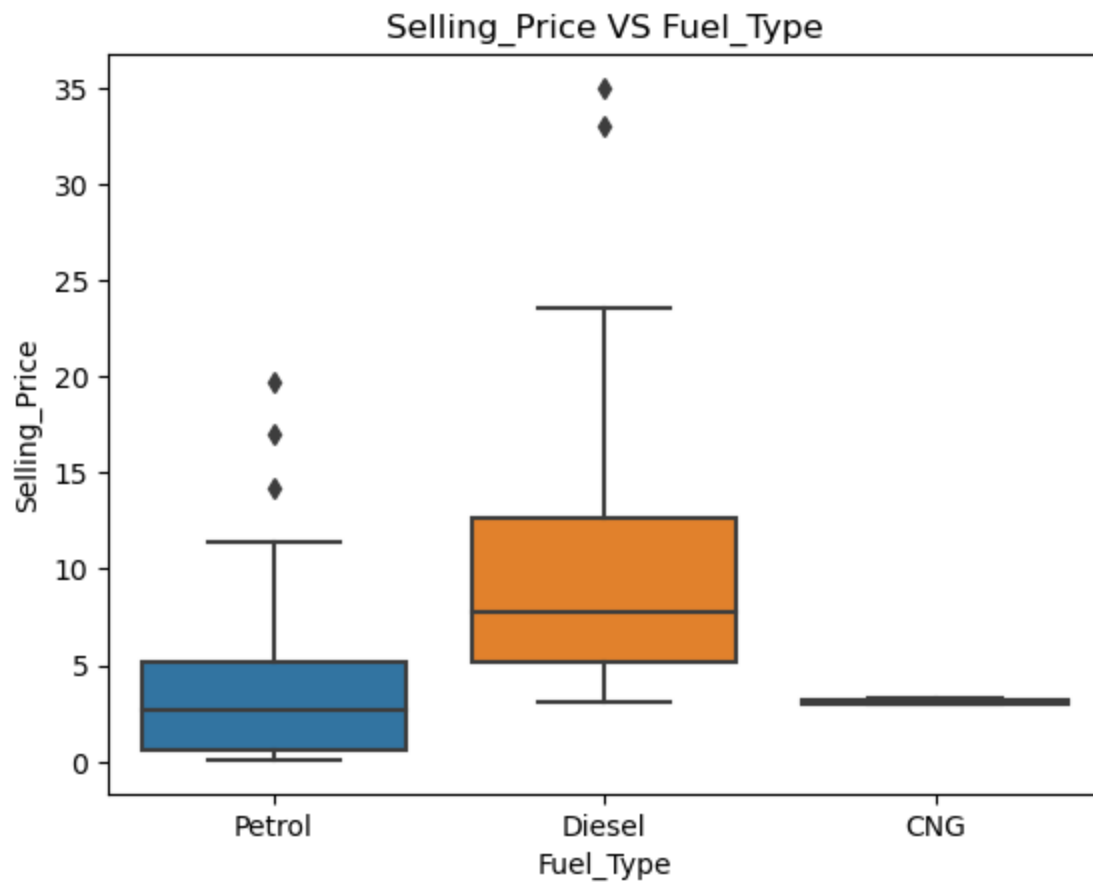
```
In [15]: 1 sns.scatterplot(x = 'Kms_Driven', y = 'Selling_Price', data = dataset).set_t
```



There appears to be no specific correlation between the selling price and kms_Driven , as observed.

Selling_Price VS Fuel_Type


```
In [16]: 1 sns.boxplot(x = 'Fuel_Type', y = 'Selling_Price', data = dataset).set_title(
```



The plot indicates that cars fueled by diesel tend to have higher selling prices. It is important to take into account the limited number of car samples with CNG fuel type in this dataset.

Selling_Price VS Seller_Type


```
In [17]: 1 sns.boxplot(x = 'Seller_Type', y = 'Selling_Price', data = dataset).set_title('Selling_Price VS Seller_Type')
```



We can observe the impact of the seller type on the selling price, which suggests that cars sold by a "Dealer" generally fetch higher selling prices.

Selling_Price VS Transmission

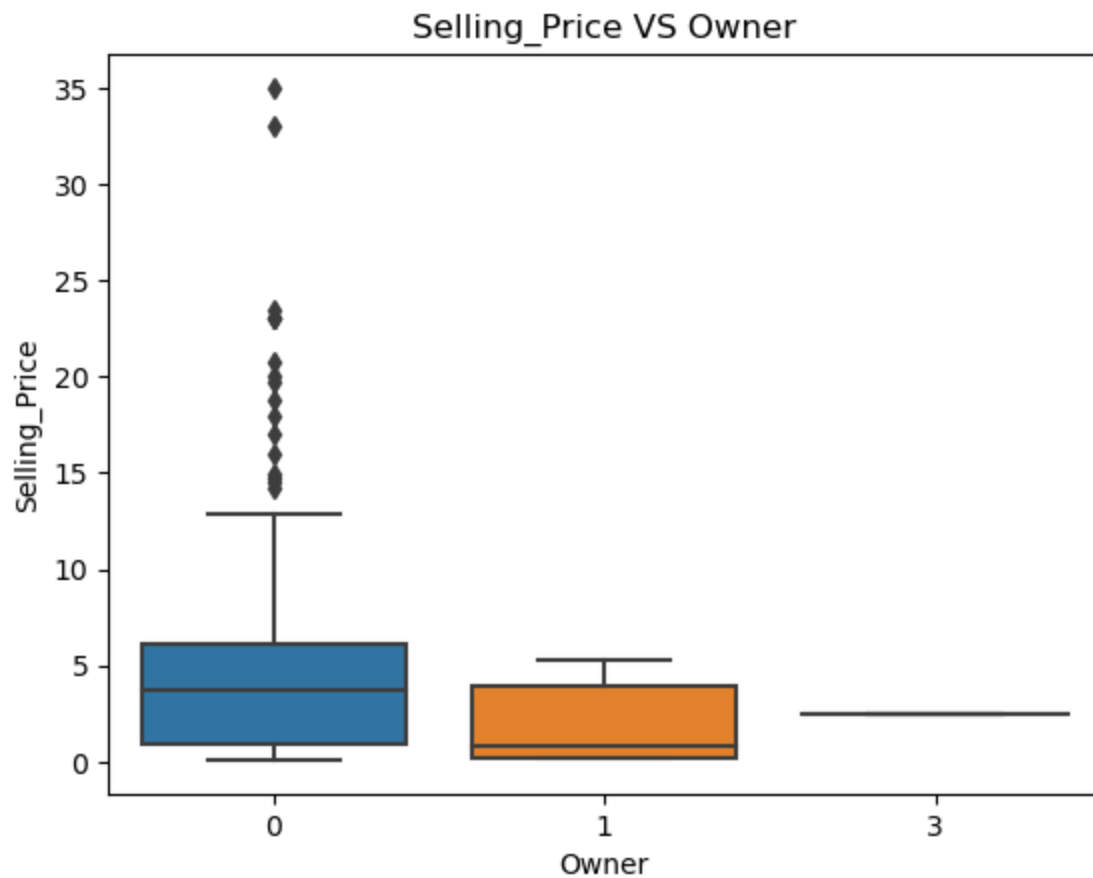
```
In [18]: 1 sns.boxplot(x = 'Transmission', y = 'Selling_Price', data = dataset).set_tit
```



The presented plot demonstrates the influence of the transmission type on the selling price. It is evident that cars equipped with automatic transmission tend to be more expensive compared to those with manual transmission. Nevertheless, there are instances where some cars with manual transmission command higher selling prices.

Selling_Price VS Owner

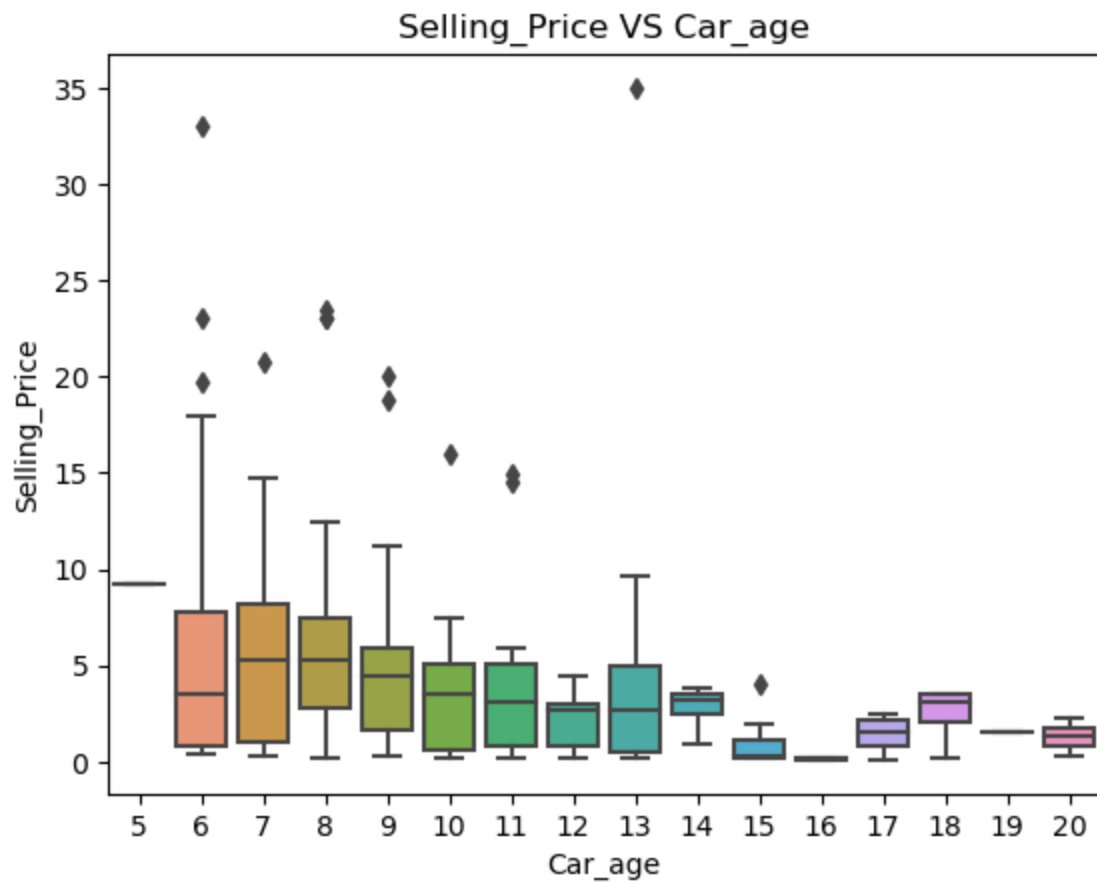

```
In [19]: 1 sns.boxplot(x = 'Owner', y = 'Selling_Price', data = dataset).set_title('Sel
```



The graph illustrates that the selling price typically declines as the number of previous owners increases.

Selling_Price VS Car_age

```
In [20]: 1 sns.boxplot(x = 'Car_age', y = 'Selling_Price', data = dataset).set_title('Selling_Price VS Car_age')
```



The last graph demonstrates that the selling price frequently decreases as the age of the car increases.

In [21]:

1 dataset

Out[21]:

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Car_a
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	
1	4.75	9.54	43000	Diesel	Dealer	Manual	0	
2	7.25	9.85	6900	Petrol	Dealer	Manual	0	
3	2.85	4.15	5200	Petrol	Dealer	Manual	0	
4	4.60	6.87	42450	Diesel	Dealer	Manual	0	
...
296	9.50	11.60	33988	Diesel	Dealer	Manual	0	
297	4.00	5.90	60000	Petrol	Dealer	Manual	0	
298	3.35	11.00	87934	Petrol	Dealer	Manual	0	
299	11.50	12.50	9000	Diesel	Dealer	Manual	0	
300	5.30	5.90	5464	Petrol	Dealer	Manual	0	

301 rows × 8 columns


In [26]:

```

1  # creat a class to plot corrilation and do onehot_encoding for given dataset
2  class manual_functions():
3      def __init__(self, dataset):
4
5          self.dataset = dataset
6
7      def corrilation(self):
8          from sklearn.preprocessing import LabelEncoder
9          global dataset
10         cat_col = [c for i, c in enumerate(dataset.columns) if dataset.dtype
11         if len(cat_col) > 0:
12             new_dataset = dataset.copy()
13             for feature in cat_col:
14                 le = LabelEncoder()
15                 label = le.fit_transform(new_dataset[feature])
16                 new_dataset.drop([feature], axis=1, inplace=True)
17                 new_dataset[feature] = label
18
19                 plt.figure(figsize = (15, 15))
20                 g= sns.heatmap(new_dataset.corr(),annot=True, )
21             else:
22                 plt.figure(figsize = (15, 15))
23                 g= sns.heatmap(new_dataset.corr(),annot=True, )
24
25         def onehot_encoding(self):
26             global dataset
27             cat_col = [c for i, c in enumerate(dataset.columns) if dataset.dtype
28             for cat_features in cat_col:
29                 dataset = pd.get_dummies(dataset, columns = [cat_features])
30             return dataset

```

In [27]: 1 manual_functions.correlation(dataset)



In this correlation matrix heatmap, we can see that the present price has highest correlation with selling price.

```
In [28]: 1 # onehot encoding
        2 manual_functions.onehot_encoding(dataset)
```

```
Out[28]:
```

	Selling_Price	Present_Price	Kms_Driven	Owner	Car_age	Fuel_Type_CNG	Fuel_Type_Diesel
0	3.35	5.59	27000	0	9	0	0
1	4.75	9.54	43000	0	10	0	1
2	7.25	9.85	6900	0	6	0	0
3	2.85	4.15	5200	0	12	0	0
4	4.60	6.87	42450	0	9	0	1
...
296	9.50	11.60	33988	0	7	0	1
297	4.00	5.90	60000	0	8	0	0
298	3.35	11.00	87934	0	14	0	0
299	11.50	12.50	9000	0	6	0	1
300	5.30	5.90	5464	0	7	0	0

301 rows × 12 columns



```
In [29]: 1 # here we need to determine the y or target
        2 target = dataset['Selling_Price']
        3 X = dataset.drop('Selling_Price', axis = 1)
        4
        5 print(X.shape, target.shape)
```

(301, 11) (301,)

Feature Importance

Feature importance is a feature selection technique usually use for larg dataset. However our dataset is a small dataset we are goting to it anyway just to show the importance of features and are not going to implement it on out model.

```
In [30]: 1 from sklearn.ensemble import ExtraTreesRegressor
        2 model = ExtraTreesRegressor()
        3 model.fit(X,target)
```

```
Out[30]: ExtraTreesRegressor()
```


In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [31]: 1 print(model.feature_importances_)
```

```
[3.23726240e-01 3.64512971e-02 4.70093518e-04 7.83290823e-02
 9.38298797e-05 2.56821706e-01 8.78417882e-03 1.36674032e-01
 5.44067201e-04 6.62484185e-02 9.18570550e-02]
```

```
In [32]: 1 feature_importance = pd.Series(model.feature_importances_, index = X.columns
2 feature_importance.nlargest(5).plot(kind = 'barh')
3 plt.show;
```



As we can see the most important feature to predict selling price is present price.

```
In [35]: 1 from sklearn.model_selection import train_test_split, cross_val_score
2 X_train, X_test, y_train, y_test = train_test_split(X, target, random_state=
3
4 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(225, 11) (76, 11) (225,) (76,)
```

Model Creation

In this notebook we are going to compare four different regressors and then tune each one that has poor performance then see which one is going to perform better.

```
In [36]: 1 import warnings
2 warnings.filterwarnings('ignore')
```

```
In [37]: 1 from sklearn.ensemble import GradientBoostingRegressor
2 from sklearn.linear_model import SGDRegressor
3 from sklearn.ensemble import RandomForestRegressor
4 from catboost import CatBoostRegressor
5
6 from sklearn.model_selection import RandomizedSearchCV
7
8 # import metrics
9 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [39]: 1 GBR_model = GradientBoostingRegressor(random_state= 42)
2 SGDR_model = SGDRegressor(random_state= 42)
3 RFR_model = RandomForestRegressor()
4 CBR_model = CatBoostRegressor()
```


GBR_model

```
In [41]: 1 GBR_model = GradientBoostingRegressor(random_state= 42)
2 GBR_model.fit(X_train, y_train)
3 y_pred = GBR_model.predict(X_test)
```

```
In [42]: 1 # evaluate model
2 GBR_MAE = mean_absolute_error(y_test, y_pred)
3 GBR_MSE = mean_squared_error(y_test, y_pred)
4 GBR_R2 = r2_score(y_test, y_pred)
5
6 # printing the results
7 print(f'MAE:{GBR_MAE}, MSE: {GBR_MSE}, R2: {GBR_R2}')
```

MAE:0.5408562819175843, MSE: 0.8808967747078672, R2: 0.9679374302804042

```
In [43]: 1 plt.scatter(y_test, y_pred);
```




As this model works pretty good with default hyperparameter we do not need to tune it

```
In [44]: 1 dft = pd.DataFrame({'Actual': y_test, 'Prediction': y_pred})  
2 dft.reset_index(drop=True, inplace=True)  
3 dft.head()
```

```
Out[44]:
```

	Actual	Prediction
0	0.35	0.424882
1	10.11	10.467838
2	4.95	4.772250
3	0.15	0.268056
4	6.95	7.657061


```
In [45]: 1 ax = sns.displot(data = (dft['Actual'], dft['Prediction']), kind = 'kde')
```



GBR model performed very good without tuning so we are not going to tune it.

SGDR_model

```
In [46]: 1 SGDR_model = SGDRegressor(random_state= 42)
2 SGDR_model.fit(X_train, y_train)
3 SGDR_y_pred = SGDR_model.predict(X_test)
```

```
In [47]: 1 SGDR_MAE = mean_absolute_error(y_test, SGDR_y_pred)
2 SGDR_MSE = mean_squared_error(y_test, SGDR_y_pred)
3 SGDR_R2 = r2_score(y_test, SGDR_y_pred)
4
5 # printing the results
6 print(f'MAE:{SGDR_MAE}, MSE: {SGDR_MSE}, R2: {SGDR_R2}')
```


```
MAE:6.4245507559862e+18, MSE: 5.884683481743032e+37, R2: -2.1418863120904382e+3
6
```

```
In [48]: 1 SGDR_df = pd.DataFrame({'Actual': y_test, 'Prediction': SGDR_y_pred})  
2 SGDR_df.head()
```

```
Out[48]:
```

	Actual	Prediction
177	0.35	-4.426537e+18
289	10.11	-2.025141e+18
228	4.95	-1.106634e+19
198	0.15	-6.455367e+18
60	6.95	-7.377746e+18

```
In [49]: 1 ax = sns.displot(data = (SGDR_df['Actual'], SGDR_df['Prediction']), kind = 'kde')
```




As we can see it cannot predict well with default hyperparameter so let's tune it.

```
In [50]: 1 loss= ['squared_error', 'huber', 'epsilon_insensitive', 'squared_epsilon_inse
2 penalty = ['l2', 'l1']
3 alpha = [float(X) for X in np.linspace(start= 0.0001, stop= 0.01, num = 100
4 l1_ratio = [float(X) for X in np.linspace(start = 0.1, stop = 0.5, num = 100
5 fit_intercept = ['True', 'False']
6 max_iter = [int(X) for X in np.linspace(start = 900, stop = 1500, num = 100)
7 learning_rate = ['constant', 'optimal', 'invscaling', 'adaptive']
8 warm_start = ['False', 'True']
```

```
In [51]: 1 random_grid = {'loss': loss,
2                       'penalty': penalty,
3                       'alpha': alpha,
4                       'l1_ratio': l1_ratio,
5                       'max_iter': max_iter,
6                       'learning_rate': learning_rate,
7                       'warm_start': warm_start,}
```

```
In [52]: 1 Tuned_SGDR_model = RandomizedSearchCV(estimator= SGDR_model, param_distribut
```



```
In [53]: 1 Tuned_SGDR_model.fit(X_train, y_train)
```

```
Out[53]: RandomizedSearchCV(cv=10, estimator=SGDRegressor(random_state=42),
                             param_distributions={'alpha': [0.0001, 0.0002,
                                                             0.00030000000000000003,
                                                             0.0004, 0.0005,
                                                             0.0006000000000000001,
                                                             0.0007000000000000001, 0.0008,
                                                             0.0009000000000000001, 0.001,
                                                             0.0011, 0.0012000000000000001,
                                                             0.0013000000000000002,
                                                             0.0014000000000000002, 0.0015,
                                                             0.0016, 0.0017000000000000001,
                                                             0.0018000000000000001, ...],
                             'loss': ['squared_error', 'huber',
                                       'epsilon_insensitive',
                                       'squared_epsilon_insensitive'],
                             'max_iter': [900, 906, 912, 918, 924,
                                           930, 936, 942, 948, 954,
                                           960, 966, 972, 978, 984,
                                           990, 996, 1003, 1009, 1016,
                                           1021, 1027, 1033, 1039,
                                           1045, 1051, 1057, 1063,
                                           1069, 1075, ...],
                             'penalty': ['l2', 'l1'],
                             'warm_start': ['False', 'True']},
                             random_state=42, scoring='neg_mean_absolute_error')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [54]: 1 Tuned_SGDR_model.best_params_
```

```
Out[54]: {'warm_start': 'False',
           'penalty': 'l1',
           'max_iter': 1324,
           'loss': 'huber',
           'learning_rate': 'adaptive',
           'l1_ratio': 0.45151515151515154,
           'alpha': 0.0089}
```

```
In [55]: 1 T_SGDR_y_pred = Tuned_SGDR_model.predict(X_test)
```

```
In [57]: 1 plt.scatter(y_test, T_SGDR_y_pred);
```



```
In [58]: 1 # evaluate model
2 T_SGDR_MAE = mean_absolute_error(y_test, T_SGDR_y_pred)
3 T_SGDR_MSE = mean_squared_error(y_test, T_SGDR_y_pred)
4 T_SGDR_R2 = r2_score(y_test, T_SGDR_y_pred)
5
6 # printing the results
7 print(f'MAE:{T_SGDR_MAE}, MSE: {T_SGDR_MSE}, R2: {T_SGDR_R2}')
```


MAE:59.07001780969827, MSE: 5217.386128116324, R2: -188.9005777859246

```
In [59]: 1 T_SGDR_df = pd.DataFrame({'Actual': y_test, 'Prediction': T_SGDR_y_pred})
2 T_SGDR_df.head()
```

```
Out[59]:
```

	Actual	Prediction
177	0.35	43.892625
289	10.11	20.069339
228	4.95	109.763463
198	0.15	64.019826
60	6.95	73.170383

```
In [60]: 1 ax = sns.displot(data = (T_SGDR_df['Actual'], T_SGDR_df['Prediction']),kind
```




Obviously, the model has not a good performance. However, it is essential to note a substantial difference between the SGDR model after the implementation of tuning techniques as compared to its pre-tuning state.

RandomForestRegressor

```
In [61]: 1 RFR_model = RandomForestRegressor()  
2 RFR_model.fit(X_train, y_train)  
3 RFR_y_pred = RFR_model.predict(X_test)
```

In [62]: 1 plt.scatter(y_test, RFR_y_pred);



In [63]: 1 RFR_MAE = mean_absolute_error(y_test, RFR_y_pred)
2 RFR_MSE = mean_squared_error(y_test, RFR_y_pred)
3 RFR_R2 = r2_score(y_test, RFR_y_pred)
4
5 print(f'MAE: {RFR_MAE} | MSE: {RFR_MSE} | R2: {RFR_R2}')


MAE: 0.5749657894736846 | MSE: 0.7935562621052629 | R2: 0.9711164194140571

In [64]: 1 RFR_df = pd.DataFrame({'Actual': y_test, 'Prediction': RFR_y_pred})
2 RFR_df.head()

Out[64]:

	Actual	Prediction
177	0.35	0.4381
289	10.11	10.8830
228	4.95	5.1575
198	0.15	0.1967
60	6.95	7.4735

```
In [65]: 1 ax = sns.displot(data = (RFR_df['Actual'], RFR_df['Prediction']), kind = 'kd
```



The random forest regressor exhibits a remarkable level of performance, rendering any further adjustments or tuning is unnecessary.


CatBoostRegressor


```
In [66]: 1 CBR_model = CatBoostRegressor()  
2 CBR_model.fit(X_train, y_train)  
3 CBR_y_pred = CBR_model.predict(X_test)
```

Learning rate set to 0.032346

0:	learn: 4.9267607	total: 141ms	remaining: 2m 21s
1:	learn: 4.8379801	total: 142ms	remaining: 1m 10s
2:	learn: 4.7517828	total: 143ms	remaining: 47.6s
3:	learn: 4.6718790	total: 144ms	remaining: 35.9s
4:	learn: 4.5993245	total: 145ms	remaining: 28.9s
5:	learn: 4.5244262	total: 146ms	remaining: 24.2s
6:	learn: 4.4477143	total: 147ms	remaining: 20.8s
7:	learn: 4.3766781	total: 148ms	remaining: 18.3s
8:	learn: 4.2933881	total: 149ms	remaining: 16.4s
9:	learn: 4.2372730	total: 149ms	remaining: 14.8s
10:	learn: 4.1637480	total: 150ms	remaining: 13.5s
11:	learn: 4.1025327	total: 151ms	remaining: 12.4s
12:	learn: 4.0415710	total: 152ms	remaining: 11.5s
13:	learn: 3.9772887	total: 153ms	remaining: 10.8s
14:	learn: 3.9159160	total: 154ms	remaining: 10.1s
15:	learn: 3.8678159	total: 156ms	remaining: 9.56s
16:	learn: 3.8074929	total: 157ms	remaining: 9.06s
17:	learn: 3.7543990	total: 158ms	remaining: 8.63s

```
In [67]: 1 plt.scatter(y_test, CBR_y_pred);
```



```
In [68]: 1 # evaluate model
2 CBR_MAE = mean_absolute_error(y_test, CBR_y_pred)
3 CBR_MSE = mean_squared_error(y_test, CBR_y_pred)
4 CBR_R2 = r2_score(y_test, CBR_y_pred)
5
6 # printing the results
7 print(f'MAE:{CBR_MAE}, MSE: {CBR_MSE}, R2: {CBR_R2}')
```


MAE:0.5097258370652804, MSE: 0.6627723092087091, R2: 0.9758766475456004

```
In [69]: 1 CBR_df = pd.DataFrame({'Actual': y_test, 'Prediction': CBR_y_pred})
2 CBR_df.head()
```

Out[69]:

	Actual	Prediction
177	0.35	0.592639
289	10.11	10.833225
228	4.95	5.120798
198	0.15	0.342871
60	6.95	7.926235

```
In [70]: 1 ax = sns.displot(data = (CBR_df['Actual'], CBR_df['Prediction']),kind = 'kde')
```



Catboost regressor also is doing well so tuning is unnecessary

In [72]:

```
1 print(f'GBR_MAE:{round(GBR_MAE, 3)}, GBR_MSE: {round(GBR_MSE, 3)}, GBR_R2: {
2 print('-----')
3 print(f'SGDR_MAE:{round(SGDR_MAE, 3)}, SGDR_MSE: {round(SGDR_MSE, 3)}, SGDR_
4 print('-----')
5 print(f'Tuned_SGDR_MAE:{round(T_SGDR_MAE, 3)}, Tuned_SGDR_MSE: {round(T_SGDR
6 print('-----')
7 print(f'RFR_MAE: {round(RFR_MAE, 3)} | RFR_MSE: {round(RFR_MSE, 3)} | RFR_R2
8 print('-----')
9 print(f'CBR_MAE:{round(CBR_MAE, 3)}, CBR_MSE: {round(CBR_MSE, 3)}, CBR_R2: {
```

GBR_MAE:0.541, GBR_MSE: 0.881, GBR_R2: 0.968

SGDR_MAE:6.424550755986201e+18, SGDR_MSE: 5.884683481743032e+37, SGDR_R2: -2.14
18863120904382e+36

Tuned_SGDR_MAE:59.07, Tuned_SGDR_MSE: 5217.386, Tuned_SGDR_R2: -188.901

RFR_MAE: 0.575 | RFR_MSE: 0.794 | RFR_R2: 0.971

CBR_MAE:0.51, CBR_MSE: 0.663, CBR_R2: 0.976

Conclusion

Based on our exploratory data analysis, it was determined that the selling price of cars is predominantly influenced by factors such as the current price, fuel type, and seller type. Specifically, a higher current price tends to correspond to a higher selling price. Moreover, vehicles with diesel fuel type tend to exhibit higher price tags, while cars sold by dealers tend to have inflated prices, potentially attributable to profit margins.

In the realm of machine learning models, we conducted a comprehensive comparison analysis involving four specific regressor models: Gradient Boosting Regressor, SGD Regressor, Random Forest Regressor, and CatBoost Regressor. Subsequent evaluation revealed that the CatBoost Regressor exhibited the highest accuracy score, garnering an impressive performance level of 97.6%.